

NS3 802.11b Throughput Performance

Freek van Tienen

4094123

30-04-2016

Abstract

With the help of NS3 simulations can be performed for several networking protocols. One of these protocols is the 802.11b protocol. This protocol allows us to communicate wireless between a STA(Station) and AP(Access Point). It is being used widely, and the most commonly used frequency range of 2.4GHz is becoming very crowded. This extensive use of the same frequencies using DSSS(Direct Sequence Spread Spectrum), makes it difficult to ensure a consistent performance in the 802.11b protocol. The amount of STA's using a single AP greatly affects the throughput of the STA's and therefor the performance. In this document we will analyse the performance impact of the amount of STA's on a single AP using NS3.

1 Problem description

The 802.11b protocol is a commonly used protocol for wireless communications between AP en STA's. When dealing with a lot of STA's on a single AP the overall performance and performance per STA should be stable enough to ensure fair access for every STA. To analyse these situations where a lot of STA's are using a single AP with a lot of data transmitted from the STA's to the AP several techniques can be used. One of the most commonly used techniques is the use of a simulator like NS3.

With the use of NS3 a full simulation of the 802.11b protocol can be done on both the physical and the software level. To simplify our problem we focus mainly on the software level and thus try to avoid all problems like Signal loss and hidden/exposed nodes. Thus for our problem we use a fixed signal strength (and thus quality) which is equal between all STA's. We also assume that no hidden and exposed nodes exists during communication. Next to that we assume that no one is interfering on our frequency and there is only one AP.

With all the above mentioned assumptions we want to analyse the throughput performance in the network under several conditions. We want to analyse what the amount of STA's in a network would mean in comparison to the data rates these STA's communicate to the AP. Also we want to take a look at what the influence of a change in packet size is for the total and average throughput in the network. This leads us to our main question: "What will be the effect of the throughput on a 802.11b network with a single AP in comparison to the number of STA's". In the next chapter we describe how we implement such simulation in NS3 and automate the gathering of the result in order to improve replicability.

2 Design and implementation

The basic 802.11b simulator consists of a single *c++* NS3 program, which constructs all simulation data. It has several options which change variables of the experiment, these include the amount of STA's, simulation time, data rate of the STA's, packet size, etc. During each run these variables are changed and eventually this data is then analysed.

To make sure that each of these runs generate use different randomisation a *runNumber* is also set as variable during each run. Next to the seed of the random number generator which is set to the time. But because I wanted it to be designed to run in parallel, to make the simulation time less, this wasn't enough hence the *runNumber*. The implementation can be seen in Figure 1.

```
/* Make sure the simulation is random */
time_t timev;
time(&timev);
RngSeedManager::SetSeed (timev);
RngSeedManager::SetRun (runNumber);
```

Figure 1: Randomness between simulations.

For the simulation of the 802.11b a default *YansWifi* physical and channel is used. The *RxGain* of the physical layer is chosen to be 0, to make sure the no difference exists between the different STA's. For the channels a *ConstantSpeedPropagationDelayModel* is chosen with a fixed loss of $-80dB$. This results in the same receiver characteristics for each STA, to obtain a fair result. How this is setup using NS3 can be seen in Figure 2.

```
/* Create WiFi and Internet stack */
NSLOG.INFO ("Installing _WiFi_and_Internet_stack.");
WifiHelper wifi;
wifi.SetStandard (WIFI_PHY_STANDARD_80211b); // Set to 802.11b

YansWifiPhyHelper wifiPhy = YansWifiPhyHelper::Default ();
wifiPhy.Set ("RxGain", DoubleValue (0));
wifiPhy.SetPcapDataLinkType (YansWifiPhyHelper::DLT_IEEE802_11_RADIO);

YansWifiChannelHelper wifiChannel = YansWifiChannelHelper::Default ();
wifiChannel.SetPropagationDelay ("ns3::ConstantSpeedPropagationDelayModel");
wifiChannel.AddPropagationLoss ("ns3::FixedRssLossModel", "Rss", DoubleValue
(-80));
wifiPhy.SetChannel (wifiChannel.Create ());
```

Figure 2: Setting up Physical and Channels.

Then to enable the mobility between STA's and obtain fair results without any hidden or exposed, it is chosen to set both the AP and the STA's at a fixed position. The position of the AP is placed 3 meters above the STA's and also separated 3 meters in x and y direction. All the STA's are fixed at the same position, and because the constant speed propagation delay model all the STA's should have the same receiving/transmitting characteristics. This allocation of the positions in the mobility model can be seen in Figure 3.

```

/* Setup mobility */
NS_LOG_INFO ("Defining_mobility.");
MobilityHelper mobility;
Ptr<ListPositionAllocator> positionAlloc = CreateObject<
    ListPositionAllocator> ();

positionAlloc->Add (Vector (0.0, 0.0, 3.0));
for(uint16_t i = 0; i < staCount; i++)
    positionAlloc->Add (Vector (3.0, 3.0, 0.0));

mobility.SetPositionAllocator (positionAlloc);
mobility.SetMobilityModel ("ns3::ConstantPositionMobilityModel");
mobility.Install (apNode);
mobility.Install (staNodes);

```

Figure 3: Setting up the mobility model.

To generate data in the network to measure the throughput a *OnOff* application is used. This application enables us to constantly transmit data for a certain amount of *on* time versus the amount of *off* time. To simulate a high amount of throughput through the network we set the *on* and *off* time such that it continuously transmits data without any *off* time. For each STA in the network such an application will be generated and will transmit data to the AP. Both the *DataRate* and *PacketSize* are variables for our experiment and are changed with the arguments of our simulation. This application is very easy to implement and will show us the performance of throughput very easily, because constant data is transmitted between the STA's and the AP. To make sure that no packets are lost this application is run on the TCP stack instead of an UDP stack. The implementation of this application can be seen in Figure 4.

```

// Create the STA application to send packets
OnOffHelper staOnOff ("ns3::TcpSocketFactory", staAddress);
staOnOff.SetAttribute ("OnTime", StringValue ("ns3::ConstantRandomVariable [
    Constant=1]"));
staOnOff.SetAttribute ("OffTime", StringValue ("ns3::ConstantRandomVariable [
    Constant=0]"));
staOnOff.SetAttribute ("PacketSize", UIntegerValue (packetSize));
staOnOff.SetAttribute ("DataRate", StringValue (dataRate));
staOnOff.SetAttribute ("Remote", AddressValue (apAddress));

```

Figure 4: The STA application.

Since we now have all the elements for a simulation we need to generate some statistics about the throughput (and thus the performance) of our network under various conditions. To do so a *csv* file is opened during the start of the experiment and with the help of a *flowMonitor* some statistics are generated. For each network flow the statistics are generated using the method in Figure 5. These statistics are then combined which will give us the total throughput, average throughput, variance in throughput between STA's and standard deviation of throughput between STA's. This calculation can be seen in 6. These total statistics are then saved to the *csv* file.

```

float throughput = i->second.rxBites * 8.0 / simulationTime / 1024 / 1024;
throughputSum += throughput;
throughputSumSq += powf(throughput, 2);

```

Figure 5: Per STA statistics.

```

double throughputMean = throughputSum / staCount;
double throughputVariance = (throughputSumSq - (powf(throughputSum, 2) /
    staCount)) / staCount;
double throughputStd = sqrtf(throughputVariance);

```

Figure 6: Total statistics.

To generate all these statistics for different data rates, packet sizes and STA's a simple python script is made. This script runs multiple simulations in parallel to speed up the process. The generation of these different runs is done through simple for loops which can be seen in Figure ??.

It will do 5 different runs for each set of variables, to remove outliers. It will test on various data rates and 2 packet size (512 bytes and 1024 bytes) and this will show us what the difference in packet size means in comparison with throughput. It will start with one single STA and will go up to 39 STA's and analyse the total throughput and average throughput for each of the runs and will save these in a new *csv* file.

```

data_rates = [ '5Mbps', '2Mbps', '1Mbps', '500Kbps', '100Kbps' ]  # Different
    data rates for simulation
packet_sizes = [512, 1024] # Different packet sizes for simulation

# Now lets build our test set
for packet_size in packet_sizes:
    for data_rate in data_rates:
        for stas in range(1, 40):
            for run in range(0, 5):
                addRun(run, stas, data_rate, packet_size)

```

Figure 7: Total statistics.

Then as last a new python script will be run to generate the graphs to visualise the in this paper. For both the total throughput and the average throughput also the variance and standard deviation are calculated between all the 5 different runs. But after analysing these result for the average throughput it seemed that the variances was almost 0 and thus the conclusion was made to leave out these results since they didn't add any value. Instead the maximum amount of standard deviation between different STA's was shown.

The source code for all these scripts can be found on GitHub(<https://github.com/fvantienen/wireless-practical>), including the results.

3 Results

The simulations of this 802.11b network are run on a Macbook Air from 2013 in Mac OS, with a Intel Core i7 1.7GHz processor and 8GB of RAM. Multiple simulations were run in parallel in order to speed up the process. In total 2000 simulations were run and 400 different type of simulations were run. Each run simulated 10 seconds of full traffic between STA's and AP, and in total 2 different type of packet sizes were chosen (1024 bytes and 512 bytes).




Figure 8: 5Mbps with packet size of 1024 bytes.




Figure 9: 5Mbps with packet size of 512 bytes.

In Figure 8 and Figure 9 we see the throughput of an 802.11b network with a data rate of 5Mbps per node from the AP to the STA. It shows us that there is practically no difference between a packet size of 1024 and 512 if we look at the total amount of throughput and the average amount of throughput per STA. Although for the lower packet size it seems that there is less deviation of throughput between the STA's, which could mean that they have a more stable connection. Both the 1024 and 512 sized packets stabilises into the same maximum throughput.




Figure 10: 2Mbps with packet size of 1024 bytes.

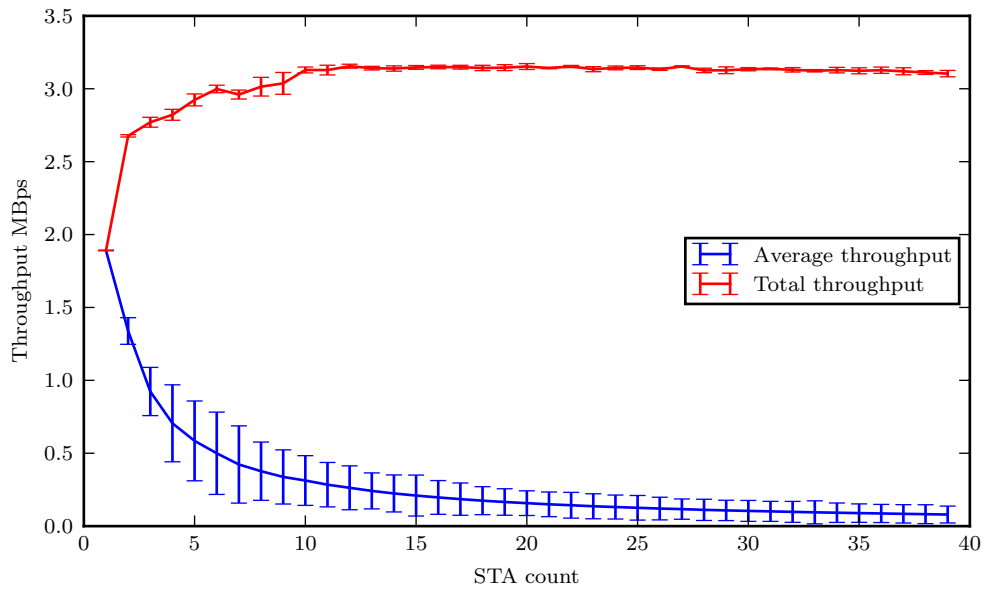


Figure 11: 2Mbps with packet size of 512 bytes.

The second run of 2Mbps shows us (in Figure 10 and Figure 11) that it takes longer to reach the stabilised total throughput, but on average the behaviour of the network is roughly the same as the 5 Mbps runs. Although you can clearly see here that there is a huge difference in deviation of the STA's between the 512 and 1024 sized packets. The 512 sized packets seem to reduce the amount of deviation in the lower STA count region (4 to 10 STA's). Also we can see that the total throughput is slowly declining if the amount of STA's increases.




Figure 12: 1Mbps with packet size of 1024 bytes.




Figure 13: 1Mbps with packet size of 512 bytes.

In the third run of 1Mbps data rate is shown in Figure 12 and Figure 13. Here you can see a specific case at the STA count of 4, where actually the deviation between STA's in the 1024 sized packets is much lower than the one from the 512 sized packets. This seems like an anomaly and can be described by the fact that these simulations are completely randomised and thus such anomalies can exists if the amount test runs isn't enough.




Figure 14: 500Kbps with packet size of 1024 bytes.




Figure 15: 500Kbps with packet size of 512 bytes.

Going even to lower data rates like 500Kbps (which is shown in Figure 14 and 15) you are able to see the stabilisation phase of the total throughput even better. In this stabilisation phase you see clearly that there is no deviation between the STA's and the maximum throughput. This can be declared by the fact that all STA's can transmit with their full data rate without any limitations in the total throughput in the network. In these graphs the effect between the 1024 and 512 sized packets cannot be seen that much anymore, because mostly the part where it had effect in the previous graphs is the part where the STA's can reach their full data rate.




Figure 16: 100Kbps with packet size of 1024 bytes.




Figure 17: 100Kbps with packet size of 512 bytes.

At the lowest data rate of 100Kbps we can see in Figure 16 and Figure 17 that the network with large amount of STA's doesn't influence the total throughput. This means that the overhead of having lot's of STA's isn't as much as having one STA with a high data rate. This can be concluded from the steepness of the curve which is exactly 1/5 of the 500Kbps curve.

4 Conclusion

To conclude I would like to say that large amount of STA's on a single AP does effect bot the total throughput and average throughput in several ways. If you look closely at the above presented graphs you can see at high data rates the total amount of throughput stays almost the same if it reached it's maximum. Although a slight overhead is present and can be seen by the slow decline in total throughput.

Next to that the lower packet sizes help to reduce the amount of deviation between STA's when the overall network hasn't reached it's total throughput yet. And since the overhead is really low, having multiple STA's with low data rate versus having one STA with a really high data rate doesn't seem to be leading into a different total throughput or average STA throughput.